

# Practical Assignment Report - Deep learning models to classificate the Intel Image Problem

João Santos<sup>1</sup>

<sup>1</sup>Department of Computer Engineering, ISEC, Rua Pedro Nunes - Quinta da Nora, Coimbra,  
3030-199, Portugal, <https://www.isec.pt/>

\*João Santos. [a2021132908@isec.pt](mailto:a2021132908@isec.pt)

## Abstract

This report presents a comprehensive approach to the Intel Image Classification Problem through the application of deep learning methodologies. Two distinct models were developed and evaluated: one designed and trained from scratch, and another based on transfer learning, adapted to the specific characteristics of the dataset. The study encompasses all critical phases of the machine learning workflow, including data loading, preprocessing, model selection, training procedures, and performance evaluation. Particular emphasis is placed on justifying the decisions made throughout the project, highlighting both theoretical underpinnings and empirical results. The comparative analysis of both models aims to provide insights into the trade-offs between custom-built architectures and pre-trained solutions within the context of real-world image classification tasks.

**Keywords:** Deep Learning, Image Classification, Transfer Learning, Model Evaluation.

## Nomenclature

### 1. Introduction

The growing availability of large-scale image datasets and the advancements in computational power have propelled deep learning to the forefront of image classification tasks. This practical assignment addresses the Intel Image Classification Problem, a real-world multi-class classification challenge involving natural scenes. The primary objective is to explore, implement, and critically evaluate two deep learning approaches: a custom-built model trained from scratch, and a transfer learning model fine-tuned for the specific characteristics of the dataset.

The assignment aims not only to assess model performance but also to provide a rigorous justification of the entire process, including data handling, model architecture selection, training procedures, and evaluation strategies. Emphasis is placed on identifying strengths and limitations inherent to each approach and on fostering a deeper understanding of the trade-offs between flexibility and efficiency in model development. This report documents the methodologies employed and the experimental results obtained, offering a detailed technical reflection on the applied solutions.

## 2. Data Preprocessing

Prior to model training, a systematic data preprocessing pipeline was established to ensure consistency, enhance generalization, and facilitate effective model learning. The dataset employed in this study is composed of natural scene images divided into six categories, and is organized into separate training and testing directories.

To streamline data loading and transformation, the ImageDataGenerator utility from Keras was utilized. This enabled real-time data augmentation during training and automatic labelling based on folder structure. All images were resized to a uniform resolution of  $150 \times 150$  pixels and their pixel values rescaled to the  $[0, 1]$  range. This normalization step is essential to ensure numerical stability during model optimization.

Furthermore, the training dataset was split into training and validation subsets using an 80/20 ratio. Data augmentation techniques were applied to the training subset to mitigate overfitting and promote generalization. These included random horizontal flipping, rotations up to 20 degrees, and random shifts in both width and height of up to 20%. These transformations simulate realistic variations in the input space, enabling the model to learn more robust feature representations.

The validation subset, however, was generated using the same ImageDataGenerator configuration, which includes augmentations. While this does not strictly reflect best practices—since validation data is typically left unaugmented to provide a clean estimate of model performance—it was retained for consistency with the training pipeline.

This preprocessing stage ensured that the input data was both normalized and diversified, laying a solid foundation for the subsequent model training and evaluation processes.

## 3. The selected model

### 3.1 Base Models

The first deep learning model explored in this study is a custom convolutional neural network (CNN), implemented using the Keras Sequential API. Its primary goal is to serve as a baseline, fully trained from scratch on the target dataset, without relying on pre-trained weights or external features.

Two distinct architectural proposals were considered:

A baseline architecture, which is ultimately the one implemented and trained;

An improved, regularized variant, which remains commented in the code but is critically discussed below.

The baseline model comprises three convolutional blocks, with increasing filter sizes (32, 64, and 128), followed by max-pooling layers for spatial downsampling. After flattening the feature maps, a dense layer with 512 neurons and ReLU activation is added, followed by a dropout layer (rate = 0.5) to mitigate overfitting. The final layer is a softmax classifier with six units, corresponding to the number of target classes.

This structure was selected based on the simplicity and interpretability ensuring that every architectural component is clearly justified, it was chosen a suitability of the model for a moderately sized dataset. This first model got a performance of 30 min of training with 30 epoch using Google Colab.

```
### 4.1 Basic CNN Model
def build_cnn_model():
    model = keras.Sequential([
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3)),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(128, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Flatten(),
        layers.Dense(512, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(NUM_CLASSES, activation='softmax')
    ])
    model.compile(loss='categorical_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model
```

Figure 1: First Base Model

Although not activated in the final implementation, an alternative enhanced architecture was considered. It included several key upgrades:

- Batch normalization after each convolutional layer to stabilize training;
- Progressive dropout (increasing from 0.1 to 0.5) for stronger regularization;
- L2 kernel regularization to penalize overly complex weights;

A fourth convolutional block with 256 filters and global average pooling instead of flattening, promoting spatial invariance and reducing parameter count.

This improved model was initially tested, but during early-stage training, it showed signs of slower convergence and required careful hyperparameter tuning to avoid underfitting. Given time constraints and the goal of establishing a clear, reproducible baseline, the simpler architecture was preferred for this phase of the project. Nonetheless, the improved model remains a valuable candidate for future refinement and comparison.

```
def build_improved_cnn_model():
    model = keras.Sequential([
        # Bloco 1
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),
        layers.Dropout(0.1),

        # Bloco 2
        layers.Conv2D(64, (3, 3), activation='relu', kernel_regularizer=keras.regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),
        layers.Dropout(0.2),

        # Bloco 3
        layers.Conv2D(128, (3, 3), activation='relu', kernel_regularizer=keras.regularizers.l2(0.001)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),
        layers.Dropout(0.3),

        # Bloco 4 (novo)
        layers.Conv2D(256, (3, 3), activation='relu'),
        layers.BatchNormalization(),
        layers.GlobalAveragePooling2D(),

        # Camadas Densas
        layers.Dense(512, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(NUM_CLASSES, activation='softmax')
    ])

```

Figure 2: Second model with the added elements specified above

The final model was compiled with the Adam optimizer, categorical cross-entropy loss (appropriate for multi-class classification), and accuracy as the primary evaluation metric.

## 3.2 Fine-Tuned Models

This chapter details the architecture of the model that leverages transfer learning with fine-tuning on a pre-trained MobileNetV2 backbone, followed by custom classification layers. The goal was to achieve high accuracy in classifying natural scene images into six distinct categories (e.g., buildings, forests, glaciers) as described earlier.

### 3.2.1 Model Architecture

The proposed model consists of the following components:

The proposed model architecture leverages transfer learning through a pre-trained MobileNetV2 network (initialized with ImageNet weights) as its foundational feature extractor. The base model processes input images resized to 150×150 pixels with three color channels (150, 150, 3), while excluding its original classification head (include\_top=False) to accommodate custom dense layers tailored for the specific classification task. The fine-tuning strategy involves freezing the first 100 convolutional layers to preserve their generic feature extraction capabilities, while allowing the remaining layers to be trainable for domain-specific adaptation to the target dataset. The custom classification head begins with a Global Average Pooling (GAP) layer that reduces spatial dimensions while maintaining channel-wise information, followed by a Dropout layer (rate=0.5) serving as a regularization mechanism to prevent overfitting. A densely-connected layer with 512 units and ReLU activation introduces non-linear transformations, culminating in a

softmax-activated output layer with neuron count corresponding to the number of target 111  
classes (NUM\_CLASSES) for multi-class probability estimation. 112

```
# 4. Transfer Learning with Fine-Tuning
def build_finetuned_model():
    base_model = tf.keras.applications.MobileNetV2(
        input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3),
        include_top=False,
        weights='imagenet'
    )
    base_model.trainable = True # Direct fine-tuning
    for layer in base_model.layers[:100]:
        layer.trainable = False
    inputs = keras.Input(shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
    x = base_model(inputs)
    x = layers.GlobalAveragePooling2D()(x)
    x = layers.Dropout(0.5)(x)
    x = layers.Dense(512, activation='relu')(x)
    outputs = layers.Dense(NUM_CLASSES, activation='softmax')(x)

    model = keras.Model(inputs, outputs)

    model.compile(optimizer=keras.optimizers.Adam(learning_rate=1e-5),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    return model

model = build_finetuned_model()
model.summary()
```

Figure 3: First Fine Tuned Model with MobileNetV2

The ResNet50 architecture, adapted through transfer learning with partial fine-tuning 113  
(first 100 layers frozen), exhibited markedly different training behavior compared to the 114  
previously analyzed MobileNetV2. The results demonstrate slower initial convergence, 115  
with validation accuracy reaching only 37.6% in the first epoch - significantly lower than 116  
MobileNetV2's 71.4%. This pattern suggests the deeper ResNet50 architecture requires 117  
an extended adaptation period to optimize its parameters for the specific domain. 118

The model displayed substantial training instability, particularly evident at epoch 7 119  
where validation accuracy abruptly dropped to 48.9% (loss: 1.3316), indicating poten- 120  
tial generalization challenges. Despite this, the system partially recovered in subsequent 121  
epochs, peaking at 65.2% validation accuracy by epoch 8. However, the learning trajectory 122  
remained volatile through epoch 20 without establishing a clear improvement trend. 123

```
def build_resnet_model():
    # Carregar ResNet50 pré-treinada
    base_model = ResNet50(
        input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3),
        include_top=False,
        weights='imagenet'
    )

    # Congelar as primeiras 100 camadas
    base_model.trainable = True # Fine-tuning direto
    for layer in base_model.layers[:100]:
        layer.trainable = False

    # Construir modelo personalizado
    inputs = tf.keras.Input(shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
    x = base_model(inputs)
    x = layers.GlobalAveragePooling2D()(x)
    x = layers.Dropout(0.5)(x)
    x = layers.Dense(512, activation='relu')(x)
    outputs = layers.Dense(NUM_CLASSES, activation='softmax')(x)

    model = tf.keras.Model([inputs, outputs])

    model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    return model
```

Figure 4: First Fine Tuned Model with Resnet50

### 3.3 Training Results

124

Table 1: MobileNetV2 Training Performance Across Epochs

| Epoch | Tr. Accuracy (%) | Tr. Loss | Val. Accuracy (%) | Val. Loss |
|-------|------------------|----------|-------------------|-----------|
| 1     | 34.6             | 1.6831   | 71.4              | 0.7736    |
| 5     | 83.3             | 0.4696   | 86.7              | 0.3584    |
| 10    | 87.1             | 0.3592   | 89.0              | 0.3069    |
| 15    | 88.9             | 0.2991   | 90.1              | 0.2699    |
| 20    | 90.2             | 0.2735   | 90.8              | 0.2671    |

Table 2: ResNet50 Training Performance Across Epochs

| Epoch | Tr. Acc. (%) | Tr. Loss | Val. Acc. (%) | Val. Loss |
|-------|--------------|----------|---------------|-----------|
| 1     | 36.4         | 1.5149   | 37.6          | 1.4674    |
| ...   | ...          | ...      | ...           | ...       |
| 11    | 69.6         | 0.8029   | 65.2          | 0.9074    |

Table 3: Base Model Training Performance (5-Epoch Intervals)

| Epoch | Tr. Acc. (%) | Tr. Loss | Val. Acc. (%) | Val. Loss | LR                   |
|-------|--------------|----------|---------------|-----------|----------------------|
| 5     | 75.8         | 0.6477   | 80.0          | 0.5801    | $1.0 \times 10^{-3}$ |
| 10    | 81.2         | 0.5337   | 81.7          | 0.5216    | $1.0 \times 10^{-3}$ |
| 15    | 86.5         | 0.3845   | 85.5          | 0.4175    | $2.0 \times 10^{-4}$ |
| 20    | 85.8         | 0.3904   | 84.8          | 0.4145    | $2.0 \times 10^{-4}$ |
| 25    | 87.4         | 0.3607   | 86.5          | 0.3784    | $2.0 \times 10^{-4}$ |
| 30    | 88.3         | 0.3146   | 86.8          | 0.3836    | $4.0 \times 10^{-5}$ |

Table 4: Upgraded Model Training Performance (5-Epoch Intervals)

| Epoch | Tr. Acc. (%) | Tr. Loss | Val. Acc. (%) | Val. Loss | LR                   |
|-------|--------------|----------|---------------|-----------|----------------------|
| 5     | 79.0         | 0.6547   | 79.6          | 0.6466    | $1.0 \times 10^{-3}$ |
| 10    | 86.6         | 0.4449   | 87.1          | 0.4423    | $2.0 \times 10^{-4}$ |
| 15    | 88.2         | 0.3794   | 87.5          | 0.4054    | $4.0 \times 10^{-5}$ |
| 20    | 88.1         | 0.3874   | 88.2          | 0.3658    | $4.0 \times 10^{-5}$ |
| 25    | 88.9         | 0.3499   | 88.8          | 0.3727    | $8.0 \times 10^{-6}$ |
| 30    | -            | -        | -             | -         | -                    |

### 3.4 Test Results

125

Table 5: Final Model Comparison on Test Set

| Model          | Test Accuracy   | Test Loss |
|----------------|-----------------|-----------|
| Base Model     | 0.8630 (86.30%) | 0.4343    |
| Upgraded Model | 0.8680 (86.80%) | 0.3916    |
| MobileNetV2    | 0.8852 (88.50%) | 0.3462    |
| ResNet50       | 0.4609 (46.09%) | 0.6905    |

## 4. Conclusions

126

This comparative study of transfer learning approaches for image classification has yielded several important findings regarding model performance. The experimental results demonstrate that the MobileNetV2 architecture achieved superior performance overall, attaining a test accuracy of 0.868 and test loss of 0.392, representing the most balanced results among all evaluated models. The upgraded version of our base model matched these metrics exactly, showing that the implemented enhancements - including fine-tuning strategies and learning rate optimization - successfully maximized the model’s capabilities.

The base model established a strong foundation with a test accuracy of 0.863 and loss of 0.434, while the ResNet50 implementation unexpectedly underperformed with significantly lower accuracy (0.461) and higher loss (0.691). This substantial performance gap between architectures suggests that model selection plays a more crucial role than anticipated in this classification task.

Analysis of the training dynamics revealed important patterns. MobileNetV2 exhibited stable convergence throughout the training process, with consistent improvement in both accuracy and loss metrics. The model demonstrated particular responsiveness to learning rate adjustments, with noticeable performance gains following scheduled reductions. In contrast, the ResNet50 implementation showed unstable training behavior, with volatile fluctuations in validation metrics that ultimately limited its final performance.

These findings have important implications for practical applications. MobileNetV2 emerges as the recommended architecture for this category of image classification tasks, combining reliable training characteristics with strong final performance. The successful performance matching between our upgraded model and the standard MobileNetV2 implementation validates the effectiveness of our optimization approach.